

Jason Sastra . Jarvis Consulting

I'm a graduate from the University of Toronto with a strong foundation in full-stack development, game programming, and data analysis. With experience across public and private sectors, and a portfolio of both technical and creative work, I'm eager to bring my skills and adaptability to a dynamic software or game development team.

Skills

Proficient: Java, Python, RDBMS/SQL, C#, C++, Agile/Scrum, Git, React.js, Power BI, Azure Web Services

Competent: Linux/Bash, Docker, IntelliJ, Pandas, Numpy, Node.js, Springboot, Kubernetes, Amazon Web Services

Familiar: Assembly, JavaFx, DBeaver, QA, Bash

Jarvis Projects

Project source code: https://github.com/jarviscanada/jarvis_data_eng_JasonSastra

Cluster Monitor [GitHub]: This project is monitoring agent for the performance and CPU usage of various different docker applications. For setup, it uses an SQL docker image to ensure SQL database could run properly. It also uses a volume to allow for continuous usage and multiple applications running. For implementation, It uses LINUX bash commands in order to obtain various statistics such as CPU usage and memory usage and periodically store them. Finally, crontab is used in order to automate the check. This project is created for users that wants to keep track of the usages of their various linux products and monitor them such that they can take action if it is underperforming or overperforming.

Core Java Apps [GitHub]: Grep App: An application that allows for the use of the grep command in Linux by utilizing Java Files Objects along with Java Input and Output Stream. It takes in a file directory and a string then searches for it on all the files for a REGEX match then output it into an output file. The project itself is managed using Maven in order to manage its dependencies and uses SLF4J Logging in order to keep track of errors. It runs by taking in its parameters as command arguments.

Python Data Analytics [GitHub]: In this project, I worked with a dataset from London Gift Shop in order to clean, process, and analyze the dataset to provide meaningful analysis as an external Data Engineer. To do so, I created a detailed Jupyter Notebook filled with RFM Segmentation Analysis, User Analysis, and Sales Analysis using Python Pandas library for data manipulation. Presented the information in a clear and concise way utilizing various graphs to showcase possible improvements to sales strategy to optimize revenue.

Databricks and Spark [GitHub]: In this project, I built a scalable data processing pipeline using Azure Databricks and PySpark to analyze large-scale financial and stock datasets, leveraging a medallion architecture (Bronze, Silver, Gold) with data stored across Azure SQL Database and Delta Lake, and integrating external APIs to generate insights visualized through Databricks dashboards. The resulting dashboard alongside the notebooks used to generate these results are additionally stored in the GitHub repository.

Springboot Trading Apps [GitHub]: In this project, I built an MVP trading application using Java and Spring Boot that is capable of retrieving real-time market quotes, process market orders, and manage trader accounts. The application follows a layered architecture consisting of controller, service, and data access layers, with Apache Tomcat handling HTTP requests and PostgreSQL serving as the database. Market data is integrated through external APIs (Finnhub and Finnage), and all endpoints are documented and testable via Swagger. I implemented RESTful controllers for quotes, traders, orders, and dashboard views, enabling account creation, deposits and withdrawals, order execution, and portfolio tracking. The system was fully containerized using Docker, with separate images for the Spring Boot application and the PostgreSQL database connected through a custom Docker network. Additionally, I wrote JUnit integration tests with a dedicated test database configuration, achieving approximately 70% test coverage.

Cloud, Kubernetes and Jenkins [GitHub]: In this project, I designed and deployed a cloud-native CI/CD pipeline using Jenkins, Docker, and Kubernetes, emphasizing application deployment across local Minikube, Azure Kubernetes Service (AKS), and a test environment on Azure VM, while automating build, containerization, and rollout processes with environment-specific pipelines and container images managed through Dockerhub. It culminated in two main environments, a local Minikube deployment which is utilized for testing and a production version in AKS.

Highlighted Projects

Into The Dataswarm [GitHub]: Developed a tactical, turn-based game with a modifiable enemy AI system that dynamically changes behavior to create engaging puzzle-solving challenges. Collaborated with a group of team until early launch on steam. • Built a modular level generator and flexible enemy patrol system, enabling seamless customization and scalability across multiple stages.

Restify: Built an AirBnB-style website for managing reservations, bookings, property listings, and user-host interactions. with a Django Backend integrated with a React Front End. The app has a responsive design that adapts seamlessly to various devices, delivering an optimized experience on both mobile and desktop.

Professional Experiences

Software Developer, Jarvis (2025/10 - present): Part of Jarvis Incubation Talent Program. Developed various Applications in Python, Java, SQL and developed in various environments including Linux, Docker and Windows as part of the training program. Integrated important technical requirements such as test driven coding in these implementations as experience.

Office of The Governing Council, University of Toronto (2023/09 - 2025/09): At the Office of the Governing Council, I managed Google Tag Manager and Google Analytics 4 to track and analyze user interactions on an internal Drupal website. I integrated GA4 data with Power BI to develop an automated, bimonthly website analytics report, and conducted Drupal testing to identify and document bugs, improving stability and user experience. Additionally, I diagnosed and resolved issues with the MERN-based OGC calendar application.

Education

University of Toronto (2020-2025), Bachelor of Arts and Science, Computer Science Specialist, Computer Science - 3.4 GPA

Miscellaneous

- Databricks Certified Data Engineer Associate (<https://credentials.databricks.com/83fe50dd-de66-45fb-a2b4-07e54b4d6d36#acc.QdfT3isx>)